

# Semaphore, Monitor, CCR,...

# Semaphores

# Semaphores

- 1965 Dijkstra
- In design of OS, we need cooperate sequential processes
- We need efficient and reliable mechanisms for supporting cooperation

# Semaphores

- ◆ Synchronization tool (provided by the OS) that do not require busy waiting
- ◆ A semaphore  $S$  is an integer variable that, apart from initialization, can only be accessed through 2 **atomic and mutually exclusive** operations:
  - $P(S)$
  - $V(S)$
- ◆ To avoid busy waiting: when a process has to wait, it will be put in a **blocked queue** of processes waiting for the same event

# Semaphores

- ◆ Hence, in fact, a semaphore is a record (structure):

```
type semaphore = record
    count: integer;
    queue: list of process
end;

var S: semaphore;
```

- ◆ When a process must wait for a semaphore S, it is blocked and put on the semaphore's queue
- ◆ The signal operation removes (according to a fair policy like FIFO) one process from the queue and puts it in the list of ready processes

# Semaphore's operations

**P(S):**

```
S.count--;  
if (S.count<0) {  
    block this process  
    place this process in S.queue  
}
```

**V(S):**

```
S.count++;  
if (S.count<=0) {  
    remove a process P from S.queue  
    place this process P on ready list  
}
```

S.count must be initialized to a nonnegative value  
(depending on application)

# Semaphores: observations

- When  $S.count \geq 0$ : the number of processes that can execute  $P(S)$  without being blocked =  $S.count$
- When  $S.count < 0$ : the number of processes waiting on  $S$  is  $= |S.count|$
- Atomicity and mutual exclusion: no 2 process can be in  $P(S)$  and  $V(S)$  (on the same  $S$ ) at the same time (even with multiple CPUs)
- Hence the blocks of code defining  $P(S)$  and  $V(S)$  are, in fact, critical sections

# Using semaphores for solving critical section problems

- For n processes
- Initialize S.count to 1
- Then only 1 process is allowed into CS (mutual exclusion)
- To allow k processes into CS, we initialize S.count to k

```
Process Pi:  
repeat  
    P(S);  
    CS  
    V(S);  
    RS  
forever
```



# Using semaphores to synchronize processes

- We have 2 processes: P1 and P2
- Statement S1 in P1 needs to be performed before statement S2 in P2
- Then define a semaphore “synch”
- Initialize synch to 0
- Proper synchronization is achieved by having in P1:  
S1;  
V(synch);
- And having in P2:  
P(synch);  
S2;


# The producer/consumer problem

- A **producer process** produces information that is consumed by a **consumer process**
  - Ex1: a print program produces characters that are consumed by a printer
  - Ex2: an assembler produces object modules that are consumed by a loader
- We need a **buffer** to hold items that are produced and eventually consumed
- A common paradigm for cooperating processes

# P/C: unbounded buffer

- We assume first an **unbounded** buffer consisting of a linear array of elements
- `in` points to the next item to be produced
- `out` points to the next item to be consumed

**shaded area indicates portion of buffer that is occupied**



# P/C: unbounded buffer

- We need a **semaphore S to perform mutual exclusion** on the buffer: only 1 process at a time can access the buffer
- We need another **semaphore N to synchronize** producer and consumer on the number  $N (= in - out)$  of items in the buffer
  - an item can be consumed only after it has been created

# Solution of P/C: unbounded buffer

Initialization:

S.count:=1;

N.count:=0;

in:=out:=0;

append(v):

b[in]:=v;

in++;

take():

w:=b[out];

out++;

return w;

Producer:

repeat

produce v;

P(S);

**append(v);**

V(S);

V(N);

forever

Consumer:

repeat

P(N);

P(S);

**w:=take();**

V(S);

consume(w);

forever

■ **critical sections**


# P/C: unbounded buffer

## • Remarks:

- Putting  $V(N)$  inside the CS of the producer (instead of outside) has no effect since the consumer must always wait for both semaphores before proceeding
- The consumer must perform  $P(N)$  before  $P(S)$ , otherwise **deadlock** occurs if consumer enter CS while the buffer is empty

• *Using semaphores is a difficult art...*

## P/C: finite circular buffer of size $k$



- can consume only when number  $N$  of (consumable) items is at least 1 (now:  $N \neq \text{in} - \text{out}$ )
- can produce only when number  $E$  of empty spaces is at least 1

# P/C: finite circular buffer of size $k$

As before:

- we need a semaphore  $S$  to have mutual exclusion on buffer access
- we need a semaphore  $N$  to synchronize producer and consumer on the number of consumable items

In addition:

- we need a semaphore  $E$  to synchronize producer and consumer on the number of empty spaces
- Producer cannot overflow the buffer



# Solution of P/C: finite circular buffer of size k

Initialization: S.count:=1; in:=0;  
N.count:=0; out:=0;  
E.count:=k;

```
append(v):  
  b[in]:=v;  
  in:=(in+1)  
    mod k;
```

```
take():  
  w:=b[out];  
  out:=(out+1)  
    mod k;  
  return w;
```


```
Producer:  
repeat  
  produce v;  
  P(E);  
  P(S);  
  append(v);
```

```
  V(S);  
  V(N);  
forever  
critical sections
```

```
Consumer:  
repeat  
  P(N);  
  P(S);  
  w:=take();  
  V(S);  
  V(E);  
  consume(w);  
forever
```

# The Dining Philosophers Problem

- 5 philosophers who only eat and think
- each need to use 2 forks for eating
- we have only 5 forks
- A classical synchronization problem
- Illustrates the difficulty of allocating resources among process without deadlock and starvation



# The Dining Philosophers Problem

- Each philosopher is a process
- One semaphore per fork:
  - fork: array[0..4] of semaphores
  - Initialization: fork[i].count:=1 for i:=0..4
- A first attempt:
- Deadlock if each philosopher start by picking his left fork!

**Process  $P_i$ :**

```
repeat  
  think;  
  P(fork[i]);  
  P(fork[i+1 mod 5]);  
  eat;  
  V(fork[i+1 mod 5]);  
  V(fork[i]);  
forever
```

# The Dining Philosophers Problem

- A solution: admit only 4 philosophers at a time that tries to eat
- Then 1 philosopher can always eat when the other 3 are holding 1 fork
- Hence, we can use another semaphore T that would limit at 4 the numb. of philosophers “sitting at the table”
- Initialize: T.count:=4


```
Process Pi:  
repeat  
    think;  
    P(T);  
    P(fork[i]);  
    P(fork[i+1 mod 5]);  
    eat;  
    V(fork[i+1 mod 5]);  
    V(fork[i]);  
    V(T);  
forever
```

# Readers/Writers Problem

- Multiple processes wanting to **read** an item, and one or more needing to **write**
  - (Think of airline reservations...)
- Rather than enforce mutual exclusion on **every** access, use these rules:
  - Any number of readers may simultaneously read the file
  - Only one writer at a time may write to the file
  - If a writer is writing to the file, no reader may read it


# Reader priority

- “First readers/writers” problem (reader priority):
  - No reader will wait (for other readers to finish) even if a writer is waiting
- Rule 1: If writer is writing, nobody can read or write
- Conditions:
  - C1: no writer is writing (for writers)
  - C2: nobody is reading or writing (for readers)
- Writer starvation possible



## Fair solution - Writer priority

- “Second reader/writers” problem (writer priority):
  - No new readers allowed once a writer has asked for access
- Rule 2: If readers are reading and at least one writers is waiting, new readers should wait
- Rule3: If writer finishes writing, all waiting readers can start read
- Three queues
- Conditions:
  - C1: no writer is writing
  - C2': nobody is reading and R is empty
  - C3: writer is writing or W is empty



## Semaphore Solution to first Readers/Writers problem (readers priority)

- Keep a count of number of current readers (readcount = 0)
- Use two semaphores
  - 1. mutex: mutually exclusive access to readcount (initially 1)
  - 2. wrt: mutual exclusion for writers (initially 1)

### Reader

```
P(mutex);  
    readcount = readcount + 1;  
    if (readcount == 1) then P(wrt);  
V(mutex);  
readTheFile();  
P(mutex);  
    readcount = readcount - 1;  
    if (readcount == 0) then V(wrt);  
V(mutex);
```

### Writer

```
P(wrt);  
writeTheFile();  
V(wrt);
```



# Binary semaphores

- The semaphores we have studied are called counting (or integer) semaphores
- We have also **binary semaphores**
  - similar to counting semaphores except that “count” is Boolean valued
  - counting semaphores can be implemented by binary semaphores...
  - generally more difficult to use than counting semaphores (eg: they cannot be initialized to an integer  $k > 1$ )

# Binary semaphores

Pb(S):

```
    if (S.value = 1) {  
        S.value := 0;  
    } else {  
        block this process  
        place this process in S.queue  
    }
```

Vb(S):

```
    if (S.queue is empty) {  
        S.value := 1;  
    } else {  
        remove a process P from S.queue  
        place this process P on ready list  
    }
```

- Semaphores provide a powerful tool for enforcing mutual exclusion and coordinate processes
- But  $P(S)$  and  $V(S)$  are scattered among several processes. Hence, difficult to understand their effects
- Usage must be correct in all the processes
- One bad (or malicious) process can fail the entire collection of processes

**CR, CCR**

# Critical Regions

- Declare a shared variable  $v$  of type  $T$  as
  - **var  $v$ : shared  $T$ ;**
    - The variable  $v$  can *only* be accessed inside a `region` statement.
  - **region  $v$  do  $S$ ;**
    - While  $S$  is being executed, no other process can access  $v$
- Using a Critical Region For Mutual Exclusion

```
var count: shared integer;
```

```
producer {  
    .  
    .  
    region count do  
        count = count + 1;  
    .  
    .  
}
```

```
consumer {  
    .  
    .  
    region count do  
        count = count - 1;  
    .  
    .  
}
```

# Implementation of a Critical Region

- Assign each shared variable,  $x$ , a semaphore,  $xmutex$ , initialized to 1
- P & V semaphore around any *region* referencing the variable
- *The Code:*

- `var x: shared T;`  
`region x do S;`

- *Becomes:*

- `var x: T;`  
`var xmutex: semaphore;`

- `P(xmutex);`

- `S;`

- `V(xmutex);`

# Possible problem

- Process 1 performs : "region x do region y do s1;"
- Process 2 performs : "region y do region x do s2;"

Process1:

```
"P(x.mutex); P(y.mutex);"  
"S1;"  
"V(y.mutex); V(x.mutex);"
```

Process2:

```
"P(y.mutex); P(x.mutex);"  
"S2;"  
"V(x.mutex); V(y.mutex);"
```

# Solution: Conditional Critical Region (CCR)

## – **region v when B do S;**

» B is a boolean expression which must be true before S is executed

**Example:** Bounded Buffer

```
var buffer: shared record
    pool: array[0..n-1] of item;
    count, in, out: integer;
end;
```

Producer:

```
region buffer when count < n
do begin
    pool[in] := nextp;
    in := in + 1 mod n;
    count := count + 1;
end;
```

Consumer:

```
region buffer when count > 0
do begin
    nextc := pool[out];
    out := out + 1 mod n;
    count := count - 1;
end;
```



# Implementation of the CCR

```
region v when B do S
var    e: semaphore; {entry semaphore}
       b: array[1..m] of sema; {delay semaphore}
           { for every condition in program }
       d: array[1..m] of int; {delay counters}
           { for every shared variable }

CCR:   P(e);
       if not Bi then
           d[i]++;
           V(e);
           P(b[i]);
       end if
       S;

       if ((B1) and (d[1] > 0)) then {d[1]--; V(b[1]); }
       elseif ((B2) and (d[2] > 0)) then {d[2]--; V(b[2]); }
       ...
       else V(e);
```

## – Problems

- large number of semaphores
- re-evaluation of all conditions
- every process must evaluate conditions of all other processes (locals!)

# Monitors

# Monitors

- Are high-level language constructs that provide equivalent functionality to that of semaphores **but are easier to control**
- Found in many concurrent programming languages
  - Concurrent Pascal, Modula-3, C++, Java...
- Can be implemented by semaphores...

# Monitor

Is a software module containing:

- One or more procedures
- An initialization sequence
- Local data variables

Characteristics:

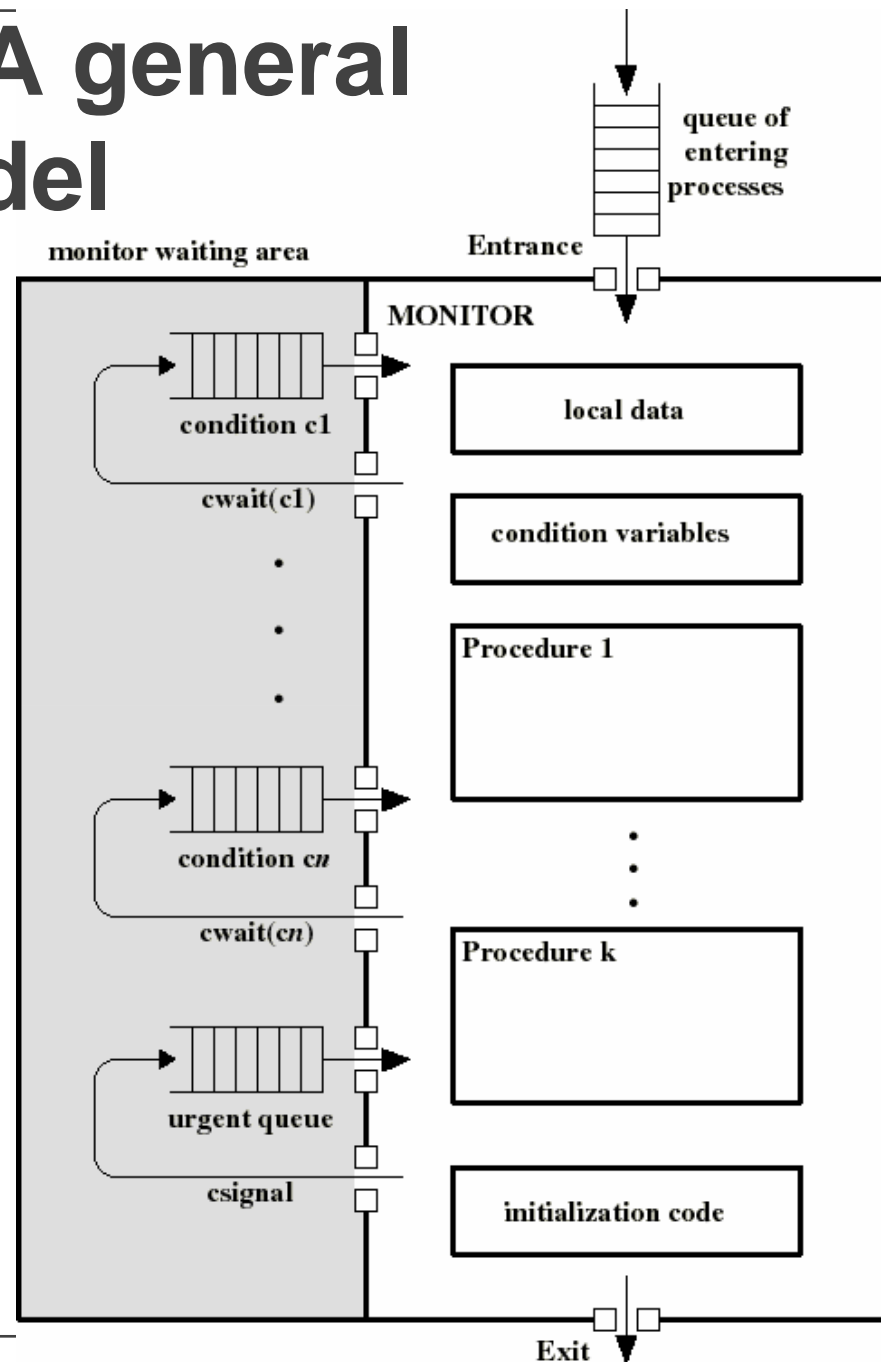
- Local variables accessible only by monitor's procedures
- A process enters the monitor by invoking one of its procedures
- Only one process can be in the monitor at any one time

# Monitor

- The monitor ensures mutual exclusion: no need to program this constraint explicitly
- Hence, shared data are protected by placing them in the monitor
  - The monitor **locks** the shared data on process entry
- Process synchronization is done by the programmer by using **condition variables** that represent conditions a process may need to wait for before executing in the monitor

# Monitor – A general model

- Awaiting processes are either in the entrance queue or in a condition queue
- A process puts itself into condition queue  $cn$  by issuing  $wait(cn)$
- $signal(cn)$  brings into the monitor 1 process in condition  $cn$  queue
- Hence  $signal(cn)$  blocks the calling process and puts it in the urgent queue (unless  $signal$  is the last operation of the monitor procedure)



# Condition variables

- Are local to the monitor (accessible only within the monitor)
- Can be access and changed only by two functions:
  - **wait(a)**: blocks execution of the calling process on condition (variable) a
    - the process can resume execution only if another process executes csignal(a)
  - **signal(a)**: resume execution of some process blocked on condition (variable) a.
    - If several such process exists: choose any one
    - If no such process exists: do nothing
  - Optional:
    - bool empty (c)
      - Returns true if delay queue is empty
    - signal\_all (c)
      - while not empty (c) do signal (c);

# Implementing monitor using semaphore

```
shared var
    e: semaphore = 1;
    c: semaphore = 0;           // for each condition
    nc: integer = 0;           // for each condition

monitor entry
    P(e)
wait(cond)
    nc = nc + 1;
    V(e); P(c); P(e);
signal(cond)
    if (nc > 0) then nc = nc - 1; V(c);
    end if
monitor exit
    V(e)
```



# Implementing semaphore using monitor

```
monitor Semaphore
  var      s: integer = 1;
          pos: condition;
  procedure P();
  begin
    while (s=0) do wait (pos);
    s = s-1;
  end;
  procedure V();
  begin
    s = s+1;
    signal(pos);
  end;
end;
```

# Monitor for the bounded P/C problem

- Monitor needs to hold the buffer:
  - buffer: array[0..k-1] of items;
- Needs two condition variables:
  - notfull: signal(notfull) indicates that the buffer is not full
  - notempty: signal(notempty) indicates that the buffer is not empty
- Needs buffer pointers and counts:
  - nextin: points to next item to be appended
  - nextout: points to next item to be taken
  - count: holds the number of items in buffer

# Monitor for the bounded P/C problem

## - the general model

Monitor boundedbuffer:

```
buffer: array[0..k-1] of items;  
nextin:=0, nextout:=0,  
    count:=0: integer;  
notfull, notempty: condition;
```

Procedure Append(v);

```
begin  
  if (count=k) wait(notfull);  
  buffer[nextin]:= v;  
  nextin:= nextin+1 mod k;  
  count++;  
  signal(notempty);  
end;
```

Procedure Take(v);

```
begin  
  if (count=0) wait(notempty);  
  v:= buffer[nextout];  
  nextout:= nextout+1 mod k;  
  count--;  
  signal(notfull);  
end;
```

# Deadlock Free Dining Philosophers

Monitor dining-philosophers

```
var state: array[0..4] of (thinking,hungry,eating);  
var philosopher: array[0..4] of condition;
```

```
procedure entry pickup( i: 0..4 );  
begin  
    state[i] := hungry;  
    test(i);  
    if state[i] != eating  
        then philosopher[i].wait;  
end;
```

```
procedure entry putdown( i: 0..4 );  
begin  
    state[i] := thinking;  
    test( i-1 mod 5 );  
    test( i+1 mod 5 );  
end;
```

# Deadlock Free Dining Philosophers

```
procedure test( k: 0..4 );  
  begin  
    if state[k-1 mod 5] != eating  
      and state[k] = hungry  
      and state[k+1 mod 5] != eating  
    then begin  
      state[k] := eating;  
      philosopher[k].signal;  
    end;  
  end;  
  
begin  
  for i := 0 to 4 do state[i] := thinking;  
end.
```

# The End